

ALMA MATER STUDIORUM · UNIVERSITÀ DI BOLOGNA

**Scuola di Scienze
Dipartimento di Fisica e Astronomia
Corso di Laurea in Fisica**

A MATHEMATICAL INTRODUCTION TO GEOMETRIC DEEP LEARNING

**Relatore:
Prof.ssa. Rita Fioresi**

**Presentata da:
Tommaso Lamma**

Anno Accademico 2020/2021

Abstract

La maggior parte delle architetture di deep learning sono costruite su modelli che imparano una partizione in classi d'equivalenza comprensibili agli umani dell'insieme delle funzioni lisce definite su domini euclidei. Sebbene questo approccio abbia successo in molte applicazioni, è limitato a una piccola classe di domini. Lo scopo del geometric deep learning è quello di estendere questo approccio a domini non euclidei come i grafi e i complessi simpliciali.

Abstract

Most of the deep learning techniques used today are based on models which learn a partition of the set of smooth functions defined on euclidean domains into human friendly equivalence classes. Although this approach has been successful in modern machine learning, it only deals with a really small set of domains. The goal of geometric deep learning is to extend this method to data defined on non euclidean domains such as graphs and simplicial complexes.

Contents

1 Preliminaries on topology	1
1.1 Simplicial Complexes	1
1.2 Abstract Simplicial Complexes	4
1.3 Simplicial Homology	5
1.4 Simplicial Cohomology	6
2 Graphs	7
2.1 Definition of Graph	7
2.2 Homology and Cohomology on Graphs	8
2.3 Laplacian Operators on Graphs	9
2.4 Laplacian Eigenfunctions in \mathbb{R}^n	11
2.5 Laplacian Eigenfunctions on Graphs	13
2.6 Heat Equation on Graphs	14
3 Geometric Deep Learning	16
3.1 Generalities on Convolution	16
3.2 The Space of Signals	17
3.3 Symmetries, Invariance and Equivariance	18
3.4 Spectral Convolution	19
A Category Theory	20
B Dirac Notation	21
C Laplacian Operators	22
Bibliography	24

Chapter 1

Preliminaries on topology

The essential idea in algebraic topology is to convert problems about topological spaces and continuous functions into problems about algebraic objects and their homomorphisms, this way one hopes to end up with an easier problem to solve.

1.1 Simplicial Complexes

In this section we shall define structures called simplicial complexes and discuss some of their properties. In order to define these structures we need the definitions of convex hull and affine independence in \mathbb{R}^n . In this chapter we recall some notions of algebraic topology, such as simplicial complexes and homology. For more details we invite the reader to consult [9], a good reference also for the preliminary necessary notions of topology we are unable to treat here.

Definition 1.1.1. Let $A \subset \mathbb{R}^n$, we define A to be *convex* if

$$x, y \in A \Rightarrow tx + (1 + t)y \in A$$

for all $t \in [0, 1]$.

In Figure 1.1 can see in blue an example of a convex set: every segment joining two points of the set lies within the set. The green set is not convex, in fact we see that the segment in the illustration partially lies outside of the set.




Fig. 1.1: Illustration of a convex (blue) and a non-convex (green) set.

Definition 1.1.2. Let $\sigma := \{x_i\}_{i \in I}$ be a subset of \mathbb{R}^n , where I is a finite set of indexes, we define σ to be *affinely independent* if $\{x_0 - x_i\}_{i \in I - \{0\}}$ is linearly independent.

We show now that the definition of affine independence of $\sigma = \{x_i\}_{i \in I} \subset \mathbb{R}^n$ is independent of the choice of x_0 .

Proposition 1.1.3. Let $\sigma := \{x_i\}_{i \in I}$ be a finite subset of \mathbb{R}^n , let $j \in I$ then, if $\{x_j - x_i\}_{i \in I - \{j\}}$ is linearly independent, also $\{x_0 - x_i\}_{i \in I - \{0\}}$ is.

Proof. If $j = 0$ the statement is trivially true. Let $j \neq 0$ and $\lambda_i \in \mathbb{R}$ for all $i \neq j$, then

$$\sum_{i \in I - \{j\}} \lambda_i (x_j - x_i) = 0 \Rightarrow \lambda_i = 0 \quad \forall i \in I - \{j\}.$$

Let then $\mu_i \in \mathbb{R}$ for all $i \neq 0$, and suppose

$$\sum_{i \in I - \{0\}} \mu_i (x_0 - x_i) = (x_0 - x_j) \sum_{i \in I - \{0\}} \mu_i + \sum_{i \in I - \{0\}} \mu_i (x_j - x_i) = 0.$$

If we define $\mu_0 := -\sum_{i \in I - \{0\}} \mu_i$ we have that

$$0 = \sum_{i \in I} \mu_i (x_j - x_i) = \sum_{i \in I - \{j\}} \mu_i (x_j - x_i) \Rightarrow \mu_i = 0 \quad \forall i \in I - \{j\},$$

which proves our proposition. the definition of affine ind the definition of affine independence is well stated. \square

Definition 1.1.4. Let $\sigma := \{x_i\}_{i \in I}$ be a finite subset of \mathbb{R}^n , we define the *convex set generated* by σ to be the smallest convex set containing X according to the inclusion relation. We shall denote this set by $[\sigma]$ and call it *convex hull* of σ .

Since the intersection of convex sets is convex, the convex set generated by σ can be equivalently defined as the intersection of all convex sets containing σ .

Theorem 1.1.5. Let $\sigma := \{x_i\}_{i \in I}$ be a finite subset of \mathbb{R}^n , if σ is affinely independent then the convex set generated by σ is

$$[\sigma] = \left\{ \sum_{i \in I} \lambda_i x_i : \lambda_i \geq 0, \sum_{i \in I} \lambda_i = 1 \right\}.$$

Furthermore for any point $x \in [\sigma]$ we have that

$$x = \sum_{i \in I} \lambda_i x_i = \sum_{i \in I} \mu_i x_i \Rightarrow \lambda_i = \mu_i \forall i \in I,$$

where $\lambda_i, \mu_i \geq 0$ and $\sum_{i \in I} \lambda_i = \sum_{i \in I} \mu_i = 1$.

Proof. Let $C := \{\cap_{\alpha} C_{\alpha} : \sigma \subset C_{\alpha}, C_{\alpha} \text{ convex}\}$, we divide the proof in three steps:

(i) $C \subset [\sigma]$.

This is true if $[\sigma]$ is convex and contains σ . The proof that it contains σ is trivial. In fact for every vertex $x_j = \sum_{i \in I} \delta_{ij} x_i$, and $\sum_{i \in I} \delta_{ij} = 1$.

To prove that it is convex we chose two points $a = \sum_{i \in I} a_i x_i, b = \sum_{i \in I} b_i x_i$ where $a_i, b_i \geq 0 \forall i \in I$ and $\sum_{i \in I} a_i = \sum_{i \in I} b_i = 1$. For $t \in [0, 1]$

$$ta + (1-t)b = t \sum_{i \in I} a_i x_i + (1-t) \sum_{i \in I} b_i x_i = \sum_{i \in I} (ta_i + (1-t)b_i) x_i.$$

Since $ta_i + (1-t)b_i \geq 0$ and $\sum_{i \in I} (ta_i + (1-t)b_i) = t \sum_{i \in I} a_i + (1-t) \sum_{i \in I} b_i = 1$ for all $i \in I$, our statement is proven.

(ii) $[\sigma] \subset C$.

If all but one the λ_i are zero certainly $\sum_{i \in I} \lambda_i x_i \in C$, since C contains all the vertexes. The inductive hypothesis, by relabeling, is that if the first $\lambda_0, \dots, \lambda_{n-1}$ are non-zero, hence

not even 1, then $\sum_{i \in I} \lambda_i x_i \in C$. We want to show that whenever $\lambda_0, \dots, \lambda_n$ are non-zero then also $\sum_{i \in I} \lambda_i x_i \in C$, since $\lambda_n \neq 1$ we have that

$$\sum_{i \in I} \lambda_i x_i = \sum_{i=0}^n \lambda_i x_i = \lambda_n x_n + \sum_{i=0}^{n-1} \lambda_i x_i = \lambda_n x_n + (1 - \lambda_n) \sum_{i=0}^{n-1} \frac{\lambda_i}{1 - \lambda_n} x_i.$$

Since $\sum_{i=0}^{n-1} \frac{\lambda_i}{1 - \lambda_n} = 1$, for the inductive hypothesis $\sum_{i=0}^{n-1} \frac{\lambda_i}{1 - \lambda_n} x_i \in C$. Also the vertex x_n is contained in C by definition, therefore, being C convex and $\lambda_n \in [0, 1]$, it follows that

$$\lambda_n x_n + (1 - \lambda_n) \sum_{i=0}^{n-1} \frac{\lambda_i}{1 - \lambda_n} x_i \in C.$$

Accordingly $\sum_{i \in I} \lambda_i x_i \in C$, by induction we conclude the proof.

(iii) $\sum_{i \in I} \lambda_i x_i = \sum_{i \in I} \mu_i x_i \Rightarrow \lambda_i = \mu_i \forall i \in I$.

Let $\sum_{i \in I} \lambda_i x_i = \sum_{i \in I} \mu_i x_i$, then also $x_0 \sum_{i \in I} \lambda_i + \sum_{i \in I} \lambda_i (x_i - x_0) = x_0 \sum_{i \in I} \mu_i + \sum_{i \in I} \mu_i (x_i - x_0)$, and since both λ_i and μ_i are normalised we have that

$$\sum_{i \in I} (\lambda_i - \mu_i)(x_0 - x_i) = \sum_{i \in I - \{0\}} (\lambda_i - \mu_i)(x_0 - x_i) = 0 \Rightarrow \lambda_i = \mu_i \quad \forall i \in I - \{0\},$$

because of the affine independence. □

Definition 1.1.6. We define a p -simplex $[\sigma]$ to be the convex hull of an affinely independent set $\sigma := \{x_i\}_{i \in I} \subset \mathbb{R}^n$, where $p = |I| - 1$ is called dimension of the p -simplex.

Theorem 1.1.5 gives us the possibility to represent a point in a simplex $[\sigma]$ via a finite set of real parameters defined in the range $[0, 1]$ and satisfying the normalisation condition $\sum_{i \in I} \lambda_i = 1$. Such parameters are called *baricentric coordinates* of $[\sigma]$.

The points in σ are called *vertexes* of the simplex $[\sigma]$, accordingly we define the vertex set of a simplex $[\sigma]$ to be $Vert([\sigma]) = \sigma$.

Definition 1.1.7. Let $[\sigma]$ be a p -simplex and $p, t \in \mathbb{N}$, we say that another t -simplex $[\tau]$ is a *face* of $[\sigma]$ or equivalently that $[\sigma]$ is a *coface* of $[\tau]$, and we write $[\tau] \leq [\sigma]$, if $\tau \subset \sigma$, where $t \leq p$.

Now we are ready for our main definitions.

Definition 1.1.8. We define a *simplicial complex* \mathcal{G} to be a collection of simplexes such that

- (i) if any simplex $[\tau] \leq [\sigma]$ and $[\sigma] \in \mathcal{G}$, then $[\tau] \in \mathcal{G}$,
- (ii) if $[\sigma], [\tau] \in \mathcal{G}$, then $[\sigma] \cap [\tau] \in \mathcal{G}$.

Figure 1.2 represents a simplicial complex, while Figure 1.3 represents a collection of simplexes which is not a simplicial complex.



Fig. 1.2: Example of simplicial complex.




Fig. 1.3: Set of simplexes which is not a simplicial complex.

In fact, we can see in Figure 1.3 that the intersection property of simplicial complexes is not satisfied.

1.2 Abstract Simplicial Complexes

In section 1.1 we studied simplicial complexes as subsets of \mathbb{R}^n . From now we shall call them *geometric simplicial complexes*. Although this approach provides simplicial complexes with the topology inherited from the metric space, it hides the power of simplicial complexes to describe networks and interactions which exist independently of that topology. To make this distinction clear, we will treat simplicial complexes as a realization of more abstract objects called *abstract simplicial complexes*. A richer discussion of abstract simplicial complexes can be found in [6] at 3.1 or in [13] at 7.3.

Definition 1.2.1. Let \mathcal{V} be a finite set, we define an *abstract simplicial complex* \mathcal{A} to be a family of non empty subsets of \mathcal{V} such that:

- (i) if $v \in \mathcal{V}$, then $\{v\} \in \mathcal{A}$,
- (ii) if $\sigma \in \mathcal{A}$ and $\tau \subset \sigma$, then $\tau \in \mathcal{A}$.

We call the members of this family *abstract simplexes*.

We call \mathcal{V} the *vertex set* of \mathcal{A} and denote it by $\text{Vert}(\mathcal{A})$; since the vertex set is finite, every abstract simplex is finite, therefore we can use the notation $\sigma = \{v_i\}_{i \in I}$, to denote a simplex in \mathcal{A} .

Definition 1.2.2. Let \mathcal{A} be an abstract simplicial complex and \mathcal{G} a geometric simplicial complex, if for all $\{x_i\}_{i \in I} \in \mathcal{A}$ also $[x_i]_{i \in I} \in \mathcal{G}$ we say that \mathcal{G} is a *geometric realization* of \mathcal{A} .

While every geometric simplicial complex can be thought as a geometric realization of an abstract simplicial complex, the existence of a geometric realization for an arbitrary abstract simplicial complex is not trivial at all.

Theorem 1.2.3. Let \mathcal{A} be an n -dimensional abstract simplicial complex, then it admits a geometric realization in \mathbb{R}^{2n+1} .

A proof of this theorem can be found in [6] at 3.1.

Both for abstract and geometric simplicial complexes one can define maps called *simplicial maps*. We obtain a category whose equivalences are called isomorphisms. A short discussion

of category theory can be found in Appendix A.

In the following sections we shall use abstract simplicial complexes, which can be always thought geometrically in the appropriate \mathbb{R}^{2d+1} .

1.3 Simplicial Homology

Homology and cohomology are key concepts in algebraic topology. We shall discuss homology theory to the extent that allows us to define the laplacian operator on simplicial complexes, for supplementary readings see [14] at 6.1 or [9] at 2.1. First we want to equip our simplicial complexes with an orientation. So far we have considered the simplex $\{x_i\}_{i \in I}$ up to reorderings of the index set I , but in most applications this is not the case.

We denote by $\{\sigma_i\}_{i \in I}$, where $I = \{1, \dots, n_p\}$, the set of all p -simplexes of the complex each one with a choice of orientation.

Definition 1.3.1. Let \mathcal{A} be a simplicial complex, we define the *group of p -chains*

$$\mathcal{C}_p(\mathcal{A}, \mathbb{R}) := \left\{ \sum_{i \in I} \lambda_i \{\sigma_i\} : \lambda_i \in \mathbb{R} \right\}.$$

In the definition 1.3.1 the sum $\sum_{i \in I} \lambda_i \{\sigma_i\}$ is the formal sum of the free abelian group on \mathbb{R} , see [11] at 1.7. This can be generalized from \mathbb{R} to an arbitrary abelian group, nevertheless, for the most of the applications, the groups $\mathbb{Z}, \mathbb{R}, \mathbb{Z}_2$ are considered.

Definition 1.3.2. Let \mathcal{A} be a simplicial complex, we define the *group of oriented p -chains* as the group $\mathcal{C}_p(\mathcal{A}, \mathbb{R})$ subject to $\{x_i\}_{i \in I} = \text{sgn}(\pi) \{x_i\}_{i \in \pi(I)}$, where $\pi : I \rightarrow I$ is any permutation of the index set. We shall denote this group as $C_p(\mathcal{A}, \mathbb{R})$, and its chains with the notation $|\sigma\rangle$.

To keep our notation light we shall write C_p instead of $C_p(\mathcal{A}, \mathbb{R})$. The constraints due to the permutations in definition 1.3.2 simply implies that whenever we swap two vertexes in the simplex we must multiply by (-1) .

Proposition 1.3.3. C_p is an abelian group.

A particularly relevant role in homology theory is played by the *boundary map*. First we define the boundary of an oriented simplex.

Definition 1.3.4. Let $|\sigma\rangle = |x_0, \dots, x_{p+1}\rangle$ be an oriented $(p+1)$ -simplex. The boundary $\partial|\sigma\rangle$ of $|\sigma\rangle$ is the p -chain defined by

$$\partial_{p+1}|\sigma\rangle := \sum_{i=0}^{p+1} (-1)^i |x_0, \dots, \widehat{x_i}, \dots, x_{p+1}\rangle$$

where the $\widehat{}$ over a symbol means that symbol is deleted.

Remark. Note that whenever we are able to construct a geometric realization for the oriented simplicial complex, the set $\bigcup_{i=0}^{p+1} [x_0, \dots, \widehat{x_i}, \dots, x_{p+1}]$ is the topological boundary of $[\sigma]$.

Furthermore we are able to extend the boundary from simplexes to chains.

Definition 1.3.5. We define the *boundary map* $\partial_{p+1} : C_{p+1} \rightarrow C_p$ to be the group homomorphism defined by

$$\partial_{p+1} \left(\sum_{i \in I} \lambda_i |\sigma_i\rangle \right) := \sum_{i \in I} \lambda_i \partial_{p+1} |\sigma_i\rangle.$$

An important property of boundary maps is the following, which we shall denote as *homology lemma*.

Lemma 1.3.6. *The boundary maps satisfy $\partial_p \circ \partial_{p+1} = 0$.*

Proof. Since the boundary maps are linear it is sufficient to check this on the generators. Let $\partial_{p+1}|x_0, \dots, x_{p+1}\rangle = \sum_{i=0}^{p+1} (-1)^i |x_0, \dots, \widehat{x}_i, \dots, x_{p+1}\rangle$ then

$$(\partial_p \circ \partial_{p+1})|x_0, \dots, x_{p+1}\rangle = \sum_{j=0, j \neq i}^{p+1} \sum_{i=0}^{p+1} (-1)^{i+j} |x_0, \dots, \widehat{x}_i, \dots, \widehat{x}_j, \dots, x_{p+1}\rangle = 0. \quad \square$$

The homology lemma is necessary to define the *homology group*. The homology group is, intuitively, the space of cycles that are not boundaries. In fact, without the homology lemma the quotient would not well defined since $\text{im } \partial_{p+1} \subset \ker \partial_p$ would not be satisfied.

Definition 1.3.7. We define the *p-homology group* to be

$$H_p := \frac{\ker \partial_p}{\text{im } \partial_{p+1}},$$

where $\text{im } \partial_{p+1}$ is called the group of simplicial *p-cycles* and $\ker \partial_p$ is called the group of simplicial *p-boundaries*.

1.4 Simplicial Cohomology

In order to define data on simplicial complexes we are interested in studying the dual of the chains $|\sigma\rangle$, which we shall call cochains $\langle\sigma|$.

Definition 1.4.1. Let C_p be the group p-chains, we define the group of *p-cochains* to be

$$C^p := \text{Hom}(C_p, \mathbb{R}).$$

Proposition 1.4.2. *C^p is an abelian group.*

Proposition 1.4.3. *The homomorphisms $\{\langle\sigma_i| : C_p \rightarrow \mathbb{R}\}$ such that $\langle\sigma_i|\sigma_j\rangle = \langle\sigma_i|(|\sigma_j\rangle) = \delta_{ij}$, form a basis of C^p .*

As well as with for chains, we have also for cochains a sequence of homeomorphisms called *coboundary maps*. The coboundary maps are defined to be the dual of the boundary maps, hence satisfying a dual version of the homology lemma called *cohomology lemma*.

Definition 1.4.4. The dual of the boundary maps which we shall call *coboundary maps*, is the group homomorphism defined by

$$d_{p+1} : C^p \rightarrow C^{p+1} \quad d_{p+1}\langle\sigma| := \langle\sigma|\partial_{p+1} \quad \forall \langle\sigma| \in C^p.$$

Therefore $d_{p+1}\langle\sigma|(|\tau\rangle) = \langle\sigma|(\partial_{p+1}|\tau\rangle) \quad \forall |\tau\rangle \in C_p$.

The proof of $d_{p+1} \circ d_p = 0$ follow directly from the homology lemma. The cohomology lemma allows us to define the *cohomology group*.

Definition 1.4.5. We define the *p-cohomology group* to be

$$H^p := \frac{\ker d_{p+1}}{\text{im } d_p},$$

where $\text{im } d_{p-1}$ is the group of simplicial p-cocycles and $\ker d_p$ is the group of simplicial p-coboundaries.

Chapter 2

Graphs

An important branch of geometric deep learning is graph representation learning. In this section we will therefore introduce graphs and graphs laplacians to build a mathematical background for graph neural networks.

2.1 Definition of Graph

Graphs are used to encode a wide variety of data, from social networks friendships to maps. We can actually consider graphs as the first topological notion introduced in mathematics. Their introduction dates back to the XVIII century. Euler noticed that in Königsberg, there was no path that allowed to cross all seven bridges just once. Hence, he started to simplify the problem to approach it in a more mathematical way.




Fig. 2.1: The city of Königsberg and the seven bridges.



Fig. 2.2: Graph representing Königsberg's seven bridges.

This led Euler to formulate an abstract approach to the problem that led to its solution, presented to the St. Petersburg Academy in 1735.

In chapter 1 we presented abstract simplicial complexes, in this chapter we shall discuss in detail 1-dimensional abstract simplicial complexes and how they are related to graphs. More information on graphs can be found in [1].

Definition 2.1.1. An *undirected graph* \mathcal{G} is an ordered pair $(V(\mathcal{G}), E(\mathcal{G}))$, consisting of a set $V(\mathcal{G})$ of *vertices* and a set $E(\mathcal{G})$, disjoint from $V(\mathcal{G})$, of *edges*, together with an incidence function $\psi_{\mathcal{G}} : E(\mathcal{G}) \rightarrow V(\mathcal{G}) \times V(\mathcal{G})$ that associates to each edge of \mathcal{G} an unordered pair of (not necessarily distinct) vertices of \mathcal{G} . If e is an edge and u and v are vertices such that $\psi_{\mathcal{G}}(e) = \{u, v\}$, then e is said to *join* u and v , and the vertices u and v are called the *ends* of e . A graph where the ordering of the ends of an edge is meaningful is said to be a *directed* graph.

Fig. 2.2 represents an undirected graph with four vertices marked in blue and seven edges.

Definition 2.1.2. Let \mathcal{G} be an undirected graph, we define the *adjacency matrix* A of \mathcal{G} as follows, A is a $n \times n$ symmetric matrix, where n is the number of vertices, such that

$$a_{ij} = \begin{cases} 1 & \text{if there exists an edge joining } i \text{ with } j \\ 0 & \text{otherwise} \end{cases}.$$

Remark. We can also define the adjacency matrix for directed graphs. For a directed graph the adjacency matrix is not symmetric.

Definition 2.1.3. A graph is *simple* if it has neither different edges connecting the same pair of vertices nor loops, i.e. there is a unique edge joining two different vertices and no vertex is joined to itself by an edge. This implies $a_{ii} = 0$ in the adjacency matrix.

Much of graph theory is concerned with the study of simple graphs.

Observation 2.1.4. We can view a *simple undirected graph* \mathcal{G} as 1-dimensional abstract simplicial complex. In fact we can construct a family of sets \mathcal{A} as in definition 1.2.1 in the following way.

$$\mathcal{A} = V(\mathcal{G}) \cup \{\{i, j\} \in V(\mathcal{G}) \times V(\mathcal{G}) : \exists e \in E(\mathcal{G}) \quad \psi_{\mathcal{G}}(e) = \{i, j\}\}.$$

It is easy to see that this is an abstract simplicial complex of dimension 1, one can readily verify the properties in def. 1.2.1.

Notation: viewing \mathcal{G} as a simplicial complex we recall that C_0 and C_1 are the groups of 0 and 1-chains respectively. We denote by $|i\rangle$ the vertex i as a group element of C_0 , and we denote $|i, j\rangle$ the edge joining $|i\rangle$ to $|j\rangle$. Since C_1 is a group, whenever $|i, j\rangle \in C_1$ then also $|j, i\rangle = -|i, j\rangle \in C_1$. A canonical basis for C_0 is the set of oriented 0-chains consisting on only one oriented 0-simplex, similarly for C_1 .

2.2 Homology and Cohomology on Graphs

Since a graph is a 1-dimensional simplicial complex, we can define only one boundary ∂_1 and one coboundary ∂_1^\dagger . We briefly recall the definitions in this context.

Definition 2.2.1. Let \mathcal{G} be a graph, we define the *boundary* $\partial_1 : C_1 \rightarrow C_0$ by

$$\partial_1 |i, j\rangle := |i\rangle - |j\rangle \quad \forall |i, j\rangle \in C_1,$$

then extended by linearity.

We define the *coboundary* $\partial_1^\dagger : C_0 \rightarrow C_1$ by

$$\partial_1^\dagger |i\rangle := \sum_{j: |i, j\rangle \in C_1} |i, j\rangle,$$

also extended by linearity.

We define the *divergence* $d_1 : C^0 \rightarrow C^1$ and the *gradient* respectively as

$$d_1 = (\partial_1)^* \quad \text{and} \quad d_1^\dagger = (\partial_1^\dagger)^*.$$

Here we identify $C_p = C^p$ with the above choice of canonical basis.

Notice that $d_1 \langle i| = \langle i| \partial_1$, where the cochain $\langle i| \partial_1$ is the dual of the chain $\partial_1^\dagger |i\rangle$.

2.3 Laplacian Operators on Graphs

On a graph that admits non trivial gradient and divergence, we can define a non trivial laplacian operator on cochains as the gradient of the divergence. Since we are going to represent it on chains we have that

$$d_1 d_1^\dagger |i\rangle = \langle i | \partial_1^\dagger \partial_1,$$

which is the adjoint of $\partial_1 \partial_1^\dagger |i\rangle$. Again we are fixing a canonical basis as above.

Definition 2.3.1. Let \mathcal{G} be a graph, we define the 0-laplacian $\Delta_0 : C_0 \rightarrow C_0$ by

$$\Delta_0 := \partial_1 \partial_1^\dagger.$$

One interesting property of the 0-laplacian, see Appendix C, is that the dimension of its kernel equals the number of connected components of the graph.

To illustrate some of the properties of graphs we will refer to the graph in Fig. 2.3.




Fig. 2.3: The graph.

$$\begin{pmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{pmatrix}.$$

Fig. 2.4: Adjacency matrix of the graph.

Example 2.3.2. Let \mathcal{G} be the graph in Figure 2.3, then the 0-laplacian expressed in terms of the canonical vertex basis $\{|i\rangle\}_{i \in I}$ is

$$\begin{pmatrix} 2 & -1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & 2 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & -1 & 2 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & -1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & -1 & 3 & -1 & -1 & 0 & 0 \\ 0 & 0 & 0 & 0 & -1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & -1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & -1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & -1 & 1 \end{pmatrix}.$$

The laplacian matrix is calculated as follows

$$\Delta_0 |1\rangle = \partial_1 |1, 2\rangle + \partial_1 |1, 3\rangle = 2 \cdot |1\rangle - |2\rangle - |3\rangle,$$

$$\Delta_0 |2\rangle = \partial_1 |2, 1\rangle + \partial_1 |2, 3\rangle = -|1\rangle + 2 \cdot |2\rangle - |3\rangle,$$

$$\Delta_0 |3\rangle = \partial_1 |3, 1\rangle + \partial_1 |3, 2\rangle = -|1\rangle - |2\rangle + 2 \cdot |3\rangle,$$

$$\Delta_0 |4\rangle = \partial_1 |4, 5\rangle = |4\rangle - |5\rangle,$$

$$\Delta_0 |5\rangle = \partial_1 |5, 4\rangle + \partial_1 |5, 6\rangle + \partial_1 |5, 7\rangle = -|4\rangle + 3 \cdot |5\rangle - |6\rangle - |7\rangle,$$

$$\Delta_0|6\rangle = \partial_1|6,5\rangle = -|5\rangle + |6\rangle,$$

$$\Delta_0|7\rangle = \partial_1|7,5\rangle = -|5\rangle + |7\rangle,$$

$$\Delta_0|8\rangle = \partial_1|8,9\rangle = |8\rangle - |9\rangle,$$

$$\Delta_0|9\rangle = \partial_1|9,8\rangle = -|8\rangle + |9\rangle.$$

Three invariant subspaces emerge from the laplacian, that determine three different laplacians, namely

$$\Delta_0 = \Delta_0^{\mathcal{A}} \oplus \Delta_0^{\mathcal{B}} \oplus \Delta_0^{\mathcal{C}}.$$

Furthermore any of those three blocks has a 1-dimensional kernel, in fact the dimensional equations for the laplacians are

$$\dim(\ker \Delta_0^{\mathcal{A}}) = \dim C_0(\mathcal{A}) - \text{rank} \begin{pmatrix} 2 & -1 & -1 \\ -1 & 2 & -1 \\ -1 & -1 & 2 \end{pmatrix} = 3 - \text{rank} \begin{pmatrix} 2 & 0 & -1 \\ 0 & 1 & -1 \\ 0 & 0 & 0 \end{pmatrix} = 3 - 2 = 1,$$

$$\dim(\ker \Delta_0^{\mathcal{B}}) = \dim C_0(\mathcal{B}) - \text{rank} \begin{pmatrix} 1 & -1 & 0 & 0 \\ -1 & 3 & -1 & -1 \\ 0 & -1 & 1 & 0 \\ 0 & -1 & 0 & 1 \end{pmatrix} = 4 - \text{rank} \begin{pmatrix} 1 & -1 & 0 & 0 \\ 0 & 2 & 0 & -1 \\ 0 & 0 & 1 & -1 \\ 0 & 0 & 0 & 0 \end{pmatrix} = 4 - 3 = 1,$$

$$\dim(\ker \Delta_0^{\mathcal{C}}) = \dim C_0(\mathcal{C}) - \text{rank} \begin{pmatrix} 1 & -1 \\ -1 & 1 \end{pmatrix} = 2 - \text{rank} \begin{pmatrix} 1 & -1 \\ 0 & 0 \end{pmatrix} = 2 - 1 = 1.$$

We can also define an higher dimensional laplacian on the graph.

Definition 2.3.3. Let \mathcal{G} be a graph, we define the 1-laplacian $\Delta_1 : C_1 \rightarrow C_1$ by

$$\Delta_1 := \partial_1^\dagger \partial_1.$$

One intersting property of the 1-laplacian, see Appendix C, is that the dimension of its kernel equals the number of independent cycles.

Example 2.3.4. In fact we can expand $\Delta_1^{\mathcal{A}} := \partial_1^\dagger \partial_1$ the basis $\{|1,2\rangle, |2,3\rangle, |3,1\rangle\}$ as

$$\begin{pmatrix} 2 & -1 & -1 \\ -1 & 2 & -1 \\ -1 & -1 & 2 \end{pmatrix}.$$

The laplacian matrix is calculated as follows

$$\Delta_1^{\mathcal{A}}|1,2\rangle = \partial_1^\dagger(|1\rangle - |2\rangle) = 2 \cdot |1,2\rangle - |3,1\rangle - |2,3\rangle,$$

$$\Delta_1^{\mathcal{A}}|2,3\rangle = \partial_1^\dagger(|2\rangle - |3\rangle) = -|1,2\rangle + 2 \cdot |2,3\rangle - |3,1\rangle,$$

$$\Delta_1^{\mathcal{A}}|3,1\rangle = \partial_1^\dagger(|3\rangle - |1\rangle) = -|1,2\rangle + 2 \cdot |3,1\rangle - |2,3\rangle.$$

We can notice that Δ_2 has a 1-dimensional kernel.

$$\dim(\ker \Delta_1^{\mathcal{A}}) = \dim C_1(\mathcal{A}) - \text{rank} \begin{pmatrix} 2 & -1 & 1 \\ -1 & 2 & 1 \\ 1 & 1 & 2 \end{pmatrix} = 3 - \text{rank} \begin{pmatrix} 2 & 0 & -1 \\ 0 & 1 & -1 \\ 0 & 0 & 0 \end{pmatrix} = 3 - 2 = 1.$$

Since the sum of the three rows is 0 we can say that the only linearly independent 1-cycle is $|1,2\rangle + |2,3\rangle + |3,1\rangle$.

2.4 Laplacian Eigenfunctions in \mathbb{R}^n

In this section we will present some properties of the eigenfunctions of the previously defined laplacian operator. The laplacian eigenfunction are the solution of a constrained variational problem for an action called *Dirichlet energy*, as shown in [2].

Definition 2.4.1. Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a differentiable function with compact support, we define the *Dirichlet energy functional* to be

$$D[f] := \int_{\mathbb{R}^n} dx |\nabla f|^2.$$

Proposition 2.4.2. If $f : \mathbb{R}^n \rightarrow \mathbb{R}$ has compact support or vanishes at infinity we have that

$$D[f] := \int_{\mathbb{R}^n} dx |\nabla f|^2 = \int_{\mathbb{R}^n} dx f \Delta f.$$

Proof. By recalling the definition of laplacian in \mathbb{R}^n as $\Delta f := -\nabla \cdot \nabla f$, and we integrate by parts

$$0 = \int_{\mathbb{R}^n} dx \nabla \cdot (f \nabla f) = \int_{\mathbb{R}^n} dx \nabla f \cdot \nabla f + \int_{\mathbb{R}^n} dx f \nabla \cdot \nabla f = \int_{\mathbb{R}^n} dx \nabla f \cdot \nabla f - \int_{\mathbb{R}^n} dx f \Delta f,$$

hence the thesis. \square

The variational problem for the Dirichlet energy leads to the Euler-Lagrange equations when it is restricted to those functions which vanish on the boundary of the domain of integration, which is true in the case of functions with compact support. Further details and proofs of the following two propositions can be found in [7] at 35 and 12 respectively.

Proposition 2.4.3. Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a sufficiently regular function with compact support, and $\mathcal{L}(f, \nabla f) : \mathbb{R}^n \rightarrow \mathbb{R}$ a sufficiently regular lagrangian, then the solutions of the variational problem

$$\delta \int_{\mathbb{R}^n} dx \mathcal{L}(f, \nabla f) = 0$$

are the solutions of the differential equation

$$\nabla \cdot \frac{\partial \mathcal{L}}{\partial(\nabla f)} - \frac{\partial \mathcal{L}}{\partial f} = 0.$$

In order to solve the variational problem subject to constraints we need a variational version of the Lagrange multipliers theorem.

Proposition 2.4.4. Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a sufficiently regular function with compact support, and $\mathcal{L}(f, \nabla f) : \mathbb{R}^n \rightarrow \mathbb{R}$ a sufficiently regular lagrangian, then the solutions of the variational problem

$$\delta \int_{\mathbb{R}^n} dx \mathcal{L}(f, \nabla f) = 0$$

constrained to the functions such that

$$\int_{\mathbb{R}^n} dx f^2 = 1$$

are the solutions of the Euler-Lagrange equations for the lagrangian

$$\mathcal{L}(f, \nabla f) - \lambda f^2,$$

where λ is a Lagrange multiplier.

We now apply this to the Dirichlet energy.

Proposition 2.4.5. *The solution of the variational problem for the Dirichlet energy*

$$\delta \int_{\mathbb{R}^n} dx |\nabla f|^2 = 0$$

constrained to the functions such that

$$\int_{\mathbb{R}^n} dx f^2 = 1$$

are the eigenfunctions of the laplacian.

Proof. For this lagrangian the Euler-Lagrange equation reads as $2\Delta f - 2\lambda f = 0$, hence

$$\Delta f = \lambda f.$$

□

Obviously, the eigenfunction that truly minimizes the Dirichlet energy is the one corresponding to the lowest eigenvalue, since for the eigenfunctions $D[f] = \int_{\mathbb{R}^n} dx f \Delta f = \lambda \int_{\mathbb{R}^n} dx f^2 = \lambda$.

Another interesting property of the laplacian eigenfunctions comes from quantum mechanical considerations. In fact, the laplacian, neglecting the constants, is the hamiltonian of the free particle in the Schrödinger representation.

Proposition 2.4.6. *The functions $e^{ip \cdot x} : \mathbb{R}^n \rightarrow \mathbb{C}$, where $p \in \mathbb{R}^n$, are a generalized orthogonal basis of laplacian eigenfunctions.*

Proof. First we see that $\Delta e^{ip \cdot x} = \nabla \cdot (i p e^{ip \cdot x}) = -p^2 e^{ip \cdot x}$. Then, in a suitable space where the Fourier transform is invertible we have that $f(x) = \int_{\mathbb{R}^n} dp e^{i \langle x, p \rangle} \int_{\mathbb{R}^n} dx' e^{-i \langle x', p \rangle} f(x')$, where $\int_{\mathbb{R}^n} dp e^{i \langle x' - x, p \rangle} = \delta(x - x')$ up to constants that can be reabsorbed in the eigenfunctions, hence the generalized completeness relation (see App. B). □

In this approach the Fourier transform is nothing more than a change of generalized basis into the basis of the eigenfunctions of the laplacian.

Let's recall the definition of convolution in \mathbb{R}^n .

Definition 2.4.7. Let $f, g : \mathbb{R}^n \rightarrow \mathbb{C}$ be sufficiently regular functions we define the *convolution* of f with g as

$$(f * g)(x) = \int_{\mathbb{R}^n} dx' f(x') g(x - x').$$

The well known theorem that follows states that the Fourier transform, as a change of generalized basis, diagonalizes the convolution.

Theorem 2.4.8. *Let \mathcal{F} denote the fourier transform, then for sufficiently regular functions $f, g : \mathbb{R}^n \rightarrow \mathbb{C}$ we have that*

$$\mathcal{F}(f * g) = (\mathcal{F}f)(\mathcal{F}g).$$

This approach will allow us in Chapter 3 to define a similar convolution as an operator which is diagonalized by the laplacian eigenfunctions.

2.5 Laplacian Eigenfunctions on Graphs

We can also define a *discrete Dirichlet energy* on graphs, using the laplacian matrix written on the canonical basis.

Definition 2.5.1. Let A, Φ be to $n \times n$ matrices, we define the *discrete Dirichlet energy* to be

$$D[\Phi] = \text{tr}(\Phi^T \Delta \Phi).$$

The minimization of the discrete Dirichlet energy is not a variational problem, rather a simple minimization in Φ in which we can use the standard Lagrange multipliers theorem.

Proposition 2.5.2. *The minimum for the Dirichlet energy $\text{tr}(\Phi^T \Delta \Phi)$ constrained to the Φ such that $\Phi^T \Phi = 1$ is reached when the rows of Φ are the laplacian eigenfunctions.*

Proof. The constrained action can also be written as

$$D[\Phi] = \sum_{i,j,k} \Phi_{ik} \Delta_{ij} \Phi_{jk} - \sum_{i,j,k} \Lambda_{ij} \Phi_{kj} \Phi_{ki},$$

where Λ is a diagonal matrix with the Lagrange multipliers as eigenvalues.

We then minimize this action

$$\begin{aligned} \frac{\partial D[\Phi]}{\partial \Phi_{mn}} &= \sum_{i,j,k} \delta_{im} \delta_{nk} \Delta_{ij} \Phi_{jk} + \sum_{i,j,k} \Phi_{ik} \Delta_{ij} \delta_{jm} \delta_{nk} - \sum_{i,j,k} \Lambda_{ij} \delta_{mk} \delta_{nj} \Phi_{ki} - \sum_{i,j,k} \Lambda_{ij} \Phi_{kj} \delta_{mk} \delta_{ni} = \\ &= \sum_j \Delta_{mj} \Phi_{jn} + \sum_i \Phi_{in} \Delta_{im} - \sum_i \Lambda_{in} \Phi_{mi} - \sum_j \Lambda_{nj} \Phi_{mj} = 0, \end{aligned}$$

which using the symmetry of Λ and A can be written as

$$\begin{aligned} \sum_j \Delta_{mj} \Phi_{jn} + \sum_i \Delta_{mi} \Phi_{in} - \sum_i \Phi_{mi} \Lambda_{in} - \sum_j \Phi_{mj} \Lambda_{jn} = \\ 2\Delta\Phi - 2\Phi\Lambda = 0. \end{aligned}$$

We then obtain

$$\Delta\Phi = \Phi\Lambda,$$

which for every row of Φ is an eigenvalue equation. □




Fig. 2.5: A graphical representation of the first laplacian eigenchains.

2.6 Heat Equation on Graphs

An relevant differential equation in physics and mathematics is the so called *heat equation*

$$\frac{df}{dt} = -\Delta f.$$

Let us see what we can say about the solution of this equation on graphs.

Proposition 2.6.1. *Let $\frac{d}{dt}|\psi\rangle = -\Delta_0|\psi\rangle$ and $|\psi\rangle|_{t=0} = |\psi_0\rangle$, then*

$$|\psi\rangle = \sum_i |e_i\rangle e^{-\lambda_i t} \langle e_i | \psi_0 \rangle,$$

where $|e_i\rangle$ and λ_i are the orthonormal eigenchains and the eigenvalues of the 0-laplacian respectively, i.e. $\Delta_0|e_i\rangle = |e_i\rangle\lambda_i$.

Proof. Since Δ_0 is selfadjoint and therefore admits the expansion $\Delta_0 = \sum_i |e_i\rangle\lambda_i\langle e_i|$, where $|e_i\rangle$ and λ_i are the orthonormal eigenchains and the eigenvalues of the 0-laplacian respectively, i.e. $\Delta_0|e_i\rangle = |e_i\rangle\lambda_i$. The discrete heat equation is linear and with constant coefficients and therefore it admits a solution of the type $|\psi\rangle = e^{-\Delta_0 t}|\psi_0\rangle$. Since $[\Delta_0, e^{-\Delta_0 t}] = 0$, these two operators share the eigenchains and therefore we also have that

$$e^{-\Delta_0 t} = \sum_i |e_i\rangle e^{-\lambda_i t} \langle e_i|,$$

hence the solution. □

From this solution we can notice that for $t \rightarrow \infty$, the only terms left are those with eigenvalue 0, namely the connected components. However, because of the term $\langle e_i | \psi_0 \rangle$, only the connected components on which $|\psi_0\rangle$ is not zero everywhere remain.

Example 2.6.2. Let's for instance take the graph in Fig. 2.3. On this graph we want to solve the heat equation given the initial condition $|\psi_0\rangle = |2\rangle + |8\rangle$. Since the subgroup generated by $\{|4\rangle, |5\rangle, |6\rangle, |7\rangle\}$ has a null projection on the initial condition, being an invariant subgroup under Δ_0 and $e^{-\Delta_0 t}$, we have that those vertices do not appear in the final solution. In order to write the solution we need to diagonalize the other two blocks of the 0-laplacian. For the block of $\{|1\rangle, |2\rangle, |3\rangle\}$, we have

$$\Delta_0 \frac{|1\rangle + |2\rangle + |3\rangle}{\sqrt{3}} = 0, \Delta_0 \frac{|3\rangle - |1\rangle}{\sqrt{2}} = 3 \frac{|3\rangle - |1\rangle}{\sqrt{2}}, \Delta_0 \frac{|2\rangle - |1\rangle}{\sqrt{2}} = 3 \frac{|2\rangle - |1\rangle}{\sqrt{2}}.$$

The two dimensional eigenspace of eigenvalue 3 can be equipped with an orthonormal basis

$$\left\{ \frac{|3\rangle - |1\rangle}{\sqrt{2}}, \frac{2|2\rangle - |1\rangle - |3\rangle}{\sqrt{2}\sqrt{3}} \right\},$$

via the Gram-Schmidt algorithm. For the block of $\{|8\rangle, |9\rangle\}$, we have

$$\Delta_0 \frac{|8\rangle + |9\rangle}{\sqrt{2}} = 0, \Delta_0 \frac{|8\rangle - |9\rangle}{\sqrt{2}} = 2 \frac{|8\rangle - |9\rangle}{\sqrt{2}}.$$

Hence the solution takes the form

$$\begin{aligned} |\psi\rangle = & \frac{|1\rangle + |2\rangle + |3\rangle}{\sqrt{3}} \frac{\langle 1| + \langle 2| + \langle 3|}{\sqrt{3}} (|2\rangle + |8\rangle) + \frac{2|2\rangle - |1\rangle - |3\rangle}{\sqrt{2}\sqrt{3}} e^{-3t} \frac{2\langle 2| - \langle 1| - \langle 3|}{\sqrt{2}\sqrt{3}} (|2\rangle + |8\rangle) + \\ & \frac{|8\rangle + |9\rangle}{\sqrt{2}} \frac{\langle 8| + \langle 9|}{\sqrt{2}} (|2\rangle + |8\rangle) + \frac{|8\rangle - |9\rangle}{\sqrt{2}} e^{-2t} \frac{\langle 8| - \langle 9|}{\sqrt{2}} (|2\rangle + |8\rangle) = \end{aligned}$$

$$= \frac{|1\rangle + |2\rangle + |3\rangle}{3} + \frac{2|2\rangle - |1\rangle - |3\rangle}{3}e^{-3t} + \frac{|8\rangle + |9\rangle}{2} + \frac{|8\rangle - |9\rangle}{2}e^{-2t}.$$

In the limit

$$\lim_{t \rightarrow \infty} |\psi\rangle = \frac{|1\rangle + |2\rangle + |3\rangle}{3} + \frac{|8\rangle + |9\rangle}{2},$$

we see that the conserved quantity that initially was 1 on both those connected components is in this limit equally shared among the vertices.

Chapter 3

Geometric Deep Learning

Learning generic functions in high dimensions is a cursed estimation problem, yet since most tasks of interest are not generic, and come with essential pre-defined regularities arising from the underlying low-dimensionality and structure of the physical world. The aim of geometric deep learning is to exploit those regularities to reduce the dimensionality of the problem.

Exploiting the symmetries of a system to reduce the dimensionality of a problem is a well known technique in physics. It is not absurd to think that symmetries can also be used in deep learning to reduce the number of independent parameters to be learnt.

The most recent formulation of geometric deep learning, which I shall follow in this chapter, is the one presented in [3].

3.1 Generalities on Convolution

Convolutional neural networks or CNN's are today one of the most common deep learning architectures, especially successful at processing that has a known grid-like topology. For instance time series, or signals can be sampled in a 1D grid, while images can be thought as a 2D grid. In Chapter 2 we defined the convolution of $f : \mathbb{R}^n \rightarrow \mathbb{R}$ with $g : \mathbb{R}^n \rightarrow \mathbb{R}$ as

$$(f * g)(x) = \int_{\mathbb{R}^n} dx' f(x') g(x - x'),$$

in the literature of convolutional neural networks the function f is called the *input* of the convolutional layer, the functions g is called *filter* and the output $f * g$ is called *feature map*.

To give an interpretation of convolution we will discuss an example from [8].

Suppose we are tracking the location of a spaceship with a laser sensor. Our sensor gives a single output $x : \mathbb{R} \rightarrow \mathbb{R}^3$, where $x(t)$ is the position read by the sensor at the time t . Now suppose that our laser sensor is somewhat noisy. To obtain a less noisy measurement of the spaceship's position, we need to average several measurements. Considering that more recent measurements are more relevant, we will want this average to be a weighted average that gives more weight to recent measurements. We can do this with a normalized weighting function $w : \mathbb{R} \rightarrow \mathbb{R}$. If we apply such a weighted average operation at every moment, we obtain a new function s providing a smoothed estimate of the position of the spaceship:

$$s(t) = \int_{\mathbb{R}} da x(a) w(t - a),$$

where $w(t - a)$ is the time-shifted weighting function. Since the acquisition rate of any measurement apparatus is not infinite, we need to see this convolution as discrete, representing time as a 1D grid,

$$s_i = \sum_{j \in \mathbb{Z}} x_j w_{(i-j)},$$

where $w_{(i-j)}$ is still the weighting function shifted by an integer. Let us see for a moment $(\mathbb{Z}, +)$ as a group, since by definition it is closed under the addition this convolution is well defined. The representation of this convolution is given by an infinite circulant matrix, therefore a restriction of \mathbb{Z} into a subset I is necessary. But what happens to $(i-j)$ at the boundary of I ? The problem is that a finite interval I is not closed under the addition in this case, namely it is not a group. Standard addition is therefore not qualified to be the group operation, yet if $I = \{0, \dots, n-1\}$, we could use addition modulo n , i.e.

$$s_i = \sum_{j=0}^{n-1} x_j w_{(i-j) \bmod n},$$

which can be seen as the product of a finite circulant matrix C_{ij} with the vector x_j . This matrix being circulant has an important constraint on the number of independent components, this reflects the compact support of the filter. This might look as a pragmatic solution but it also tells us some truth about the symmetries of the task. In some cases the classification of a signal, e.g. voice recognition, is independent of the time when the signal was recorded, we can say that the classification is invariant under translations of the signal over time. This is reflected in the fact that if I shift the whole cyclic group the matrix representing the convolution will still be circulant. At this point, we could see our filter as a function defined on the cyclic group itself, rather than a filter whose symmetry group is the cyclic group. This subtle distinction allowed M. Welling and T. S. Cohen to define in [4] a group equivariant convolution on an arbitrary group.

3.2 The Space of Signals

While a linear combination of points of a domain is not necessarily defined, the space of signals on that domain can be equipped with a vector space structure. More formally we can define the space of V -valued signals on the domain Ω .

Definition 3.2.1. Let Ω be a set, possibly with additional structure¹, and V a vector space over \mathbb{K} , we define the *signal set*

$$\mathcal{S}(\Omega, V) = \{\phi : \Omega \rightarrow V\},$$

where the dimensions of V are called *channels*.

In particular cases such as when Ω is a smooth manifold we might also require the signals to be smooth. Let us now equip the signal set with a vector space structure.

Proposition 3.2.2. *The signal set equipped with the operations*

$$\cdot : \mathbb{K} \times \mathcal{S}(\Omega, V) \rightarrow \mathcal{S}(\Omega, V) \quad \text{such that } (\lambda\phi)(x) = \lambda\phi(x), \text{ and}$$

$$+ : \mathcal{S}(\Omega, V) \times \mathcal{S}(\Omega, V) \rightarrow \mathcal{S}(\Omega, V) \quad \text{such that } (\phi + \eta)(x) = \phi(x) + \eta(x),$$

where $x \in \Omega$, is a vector space over \mathbb{K} .

We shall call this vector space *signal space*, and denote it just by \mathcal{S} for a lighter notation.

Given an inner product $\langle \cdot, \cdot \rangle_V$ in V and a measure μ on Ω the signal space can be equipped with a Hilbert space structure.

Definition 3.2.3. Let $\langle \cdot, \cdot \rangle_V : V \times V \rightarrow \mathbb{K}$ be an inner product and let μ be a measure on Ω , then for any two signals $\phi, \eta \in \mathcal{S}$ we define their inner product $\langle \cdot, \cdot \rangle : \mathcal{S} \times \mathcal{S} \rightarrow \mathbb{K}$ as

$$\langle \phi, \eta \rangle := \int_{\Omega} d\mu(x) \langle \phi(x), \eta(x) \rangle_V.$$

¹For instance Ω is a smooth manifold.

Proposition 3.2.4. *The space \mathcal{S} equipped with \langle, \rangle is a Hilbert space.*

Non vorrei dimostrarlo ma in teoria non devo imporre che \mathbb{K} sia completo per avere la completezza di \mathcal{S} ?

When the domain Ω is discrete we can choose μ to be the counting measure, therefore turning the integral into a sum.

3.3 Symmetries, Invariance and Equivariance

The symmetries of an object are colloquially transformations that leave a certain *classification* of said object unchanged or invariant.

Definition 3.3.1. We define a *classification* to be a function $C : \mathcal{S} \rightarrow \mathcal{L}$ from the signal space into some set \mathcal{L} .

In many classifications task the set \mathcal{L} is a label set.

Using the concept of classification we could give a more formal definition of symmetry.

Definition 3.3.2. Let $C : \mathcal{S} \rightarrow \mathcal{L}$ be a classification we define a *symmetry* of C to be an invertible function $g : \mathcal{S} \rightarrow \mathcal{S}$ such that

$$C[g(\phi)] = C[\phi] \quad \forall \phi \in \mathcal{S}.$$

We will denote the set of these symmetries as $S_C := \{g : \mathcal{S} \rightarrow \mathcal{S} : C[g(\phi)] = C[\phi] \forall \phi \in \mathcal{S}\}$

Proposition 3.3.3. *Let \circ be the composition of functions from \mathcal{S} to itself, then (S_C, \circ) is a group.*

Proof. We shall prove that this set with respect to \circ satisfies all the group properties.

- (i) The operation \circ is already associative.
- (ii) Let $\text{id}_{\mathcal{S}} : \mathcal{S} \rightarrow \mathcal{S}$ be the identity on \mathcal{S} , namely $\text{id}_{\mathcal{S}}(\phi) = \phi$ for all $\phi \in \mathcal{S}$, then $\text{id}_{\mathcal{S}} \in S_C$.
- (iii) Let $g, g' \in S_C$, we have that $C[(g \circ g')(\phi)] = C[\phi]$, therefore $(g \circ g') \in S_C$.
- (iv) Let $g \in S_C$ then $C[g^{-1}(\phi)] = C[\phi]$, since $C[(g \circ g^{-1})(\phi)] = C[(\phi)] = C[g^{-1}(\phi)]$. □

It is important to notice that in order to define a symmetry we need to choose a classification, namely there is no a priori symmetry in the data until we specify what we want to do with it.

Another way to see how the group acts on data is to see the group itself as an abstract object and define an *action* of this group on the domain of the signals.

Definition 3.3.4. Let G be a group and Ω a set, then we define *action* of G on Ω a mapping $\cdot : G \times \Omega \rightarrow \Omega$, such that $(g, x) \mapsto g \cdot x$ and for all $g, h \in G$ $g \cdot (h \cdot x) = (g \circ h) \cdot x$, i.e. it is compatible with the group operation.

From a group action on the domain Ω we can obtain a group action on the signal space.

Proposition 3.3.5. *Let G be a group and let \mathcal{S} be the space of signals from Ω to V , then for all $\phi \in \mathcal{S}$ and for all $g \in G$ we have that*

$$(g \cdot \phi)(x) := \phi(g^{-1} \cdot x),$$

where $x \in \Omega$, is an action of G on \mathcal{S} .

Proof. We can see that $g.(h.\phi)(x) = g.\phi(h^{-1}.x) = \phi((g^{-1} \circ h^{-1}).x) = ((g \circ h).\phi)(x)$. \square

An important class of group actions is that of *linear* group actions, also called *group representations*. A linear group action has the additional property $g.(\lambda\phi + \mu\psi) = \lambda g.\phi + \mu g.\psi$, for all ϕ, ψ in the set on which the group acts. A linear group action allows us to associate to an element g of G a $n \times n$ invertible matrix.

Definition 3.3.6. Let G be a group, we define an n -dimensional representation of G to be a homomorphism $\rho : G \rightarrow GL(n)^2$.

An important role in geometric deep learning is played by the so called *group equivariant* and *group invariant* functions.

Definition 3.3.7. Let \mathcal{C} be some set, a function $f : \mathcal{S} \rightarrow \mathcal{C}$, is called G -invariant if

$$f(\rho(g)\phi) = f(\phi) \quad \forall g \in G,$$

where $\phi \in \mathcal{S}$.

Definition 3.3.8. A function $f : \mathcal{S} \rightarrow \mathcal{S}$ is called G -equivariant if

$$f(\rho(g)\phi) = \rho(g)f(\phi) \quad \forall g \in G,$$

where $\phi \in \mathcal{S}$.

3.4 Spectral Convolution

Another approach that can be used to define a convolution is the spectral convolution with respect to a laplacian operator, in this section we shall see this abstract and general method on graphs. In order to construct a convolutional neural network that classifies data defined on a graph we need a definition of convolution on graphs. To do this we recall that in Chapter 2 we saw that the Fourier tranform diagonalizes the convolution. Here we want to define a convolution on graphs based on this fact. First we introduce what we call *graph Fourier transform*.

Definition 3.4.1. Let $|f\rangle \in C_0$ and $\dim C_0 = n_0$, we define the *graph Fourier transform* $\mathcal{F}_0 : \mathbb{R}^{n_0} \rightarrow \mathbb{R}^{n_0}$ to be

$$(\langle i|f\rangle)_{i \in I} \mapsto (\langle e_i|f\rangle)_{i \in I},$$

where $I = \{1, \dots, n_0\}$ and $|e_i\rangle$ are the eigenfunctions of the 0-laplacian.

This transform only defines a change of basis since $|f\rangle = \sum_{i \in I} \langle i|f\rangle |i\rangle = \sum_{i \in I} \langle e_i|f\rangle |e_i\rangle$, and therefore is invertible. We can in fact represent the graph Fourier transform with the matrix $F_{ij}^{-1} = F_{ij}^\dagger := \langle e_i|j\rangle$, and its inverse $F_{ij} := \langle i|e_j\rangle$. To define a convolution between two 0-chains we use the famous convolution theorem $\mathcal{F}(f * \psi) = \mathcal{F}(f)\mathcal{F}(\psi)$.

Definition 3.4.2. Let $\{|e_i\rangle\}_{i \in I}$ be a basis such that $\Delta_p |e_i\rangle = \lambda_i |e_i\rangle$, let $|f\rangle, |\psi\rangle \in C_p$, we define the representatives of $|f * \psi\rangle$ on the laplacian eigenchains to be

$$\langle e_i|f * \psi\rangle := \langle e_i|f\rangle \langle e_i|\psi\rangle \quad \forall i \in I.$$

Therefore $|f * \psi\rangle = \sum_{i \in I} \langle e_i|f\rangle \langle e_i|\psi\rangle |e_i\rangle$.

This can be trivially extended to simplicial complexes³ using the eigenfunctions of higher dimensional laplacians, as done in [5].

² $GL(n)$ is the group of $n \times n$ invertible matrices.

³Actually on any chain complex.

Appendix A

Category Theory

Definition A.1. A category \mathbf{C} consists of three ingredients:

1. a class of *objects* $Obj(\mathbf{C})$,
2. sets of *morphisms* $Hom(A, B)$ for every ordered pair $(A, B) \in Obj(\mathbf{C}) \times Obj(\mathbf{C})$,
3. a composition $Hom(A, B) \times Hom(B, C) \rightarrow Hom(A, C)$, denoted by $(f, g) \mapsto f \circ g$ for every $A, B, C \in Obj(\mathbf{C})$, satisfying the following axioms:
 - (i) the family of $Hom(A, B)$ is pairwise disjoint,
 - (ii) the composition, when defined, is associative,
 - (iii) for each $A \in Obj(\mathbf{C})$ there exists an *identity* $1_A \in Hom(A, A)$ such that for $f \in Hom(A, B)$ and $g \in Hom(C, A)$ we have that $1_A \circ f = f$ and $g \circ 1_A = g$.

Instead of writing $f \in Hom(A, B)$, we usually write $f : A \rightarrow B$.

Definition A.2. Let \mathbf{A} and \mathbf{C} be categories, a *functor* $T : \mathbf{A} \rightarrow \mathbf{C}$ is a function, that is,

- (i) for each $A \in Obj(\mathbf{A})$ it assigns $TA \in Obj(\mathbf{C})$,
- (ii) for each morphism $f : A \rightarrow A'$ it assigns a morphism $Tf : TA \rightarrow TA' \quad \forall A, A' \in Obj(\mathbf{A})$,
- (iii) if f, g are morphisms in \mathbf{A} for which $g \circ f$ is defined, then $T(g \circ f) = (Tg) \circ (Tf)$,
- (iv) $T(1_A) = 1_{TA} \quad \forall A \in \mathbf{A}$.

The property (iii) of the previous definition actually defines what we shall call *covariant functors*. If instead we require $T(g \circ f) = (Tf) \circ (Tg)$, we are defining a *contravariant functor*.

Definition A.3. An *equivalence* in a category \mathbf{C} is a morphism $f : A \rightarrow B$ for all $A, B \in Obj(\mathbf{C})$ for which there exists a morphism $g : B \rightarrow A$ such that $f \circ g = 1_B$ and $g \circ f = 1_A$.

Theorem A.4. If \mathbf{A} and \mathbf{C} are categories and $T : \mathbf{A} \rightarrow \mathbf{C}$ is a functor of either variance, then whenever f is an equivalence on \mathbf{A} then Tf is an equivalence on \mathbf{C} .

Proof. We apply T to the equations $f \circ g = 1_B$ and $g \circ f = 1_A$, that for a covariant functor leads to $(Tf) \circ (Tg) = T(1_B) = 1_{TB}$ and $(Tg) \circ (Tf) = T(1_A) = 1_{TA}$. \square

A category that will be used in the following section is the category of topological spaces and continuous functions.

Proposition A.5. Topological spaces and continuous functions are a category **Top**, whose equivalences are called homeomorphisms.

Other examples of categories can be found in [13] at 0.3 and in [12].

Appendix B

Dirac Notation

In the Dirac notation we use the kets $|\psi\rangle$ as vectors (group elements) and the bras $\langle\psi|$ as their dual. The bra-ket product $\langle\psi|\varphi\rangle = \langle\psi|(\varphi)$ is the bra evaluated on the ket, this is an inner product of $|\psi\rangle$ and $|\varphi\rangle$. In this notation the spectral theorem for selfadjoint operators with a discrete spectrum reads as

$$A = \sum_i |a_i\rangle a_i \langle a_i|,$$

where $A|a_i\rangle = a_i|a_i\rangle$, $\langle a_i|a_j\rangle = \delta_{ij}$ and $\sum_i |a_i\rangle\langle a_i| = 1$. The relation $\sum_i |a_i\rangle\langle a_i| = 1$ is called completeness relation and is valid only if $\{|a_i\rangle\}$ is an orthonormal basis, 1 is the identity operator. An extension of the concept of orthonormal basis is the concept of generalized orthonormal basis. Generalized orthonormal eigenfunctions allows us to diagonalize an operator with a continuous spectrum as follows

$$A = \int da |a\rangle a \langle a|,$$

where $A|a\rangle = a|a\rangle$, $\langle a|a'\rangle = \delta(a - a')$ and $\int da |a\rangle\langle a| = 1$. The relation $\int da |a\rangle\langle a| = 1$ is called generalized completeness relation and is valid only if $\{|a\rangle\}$ is a generalized orthonormal basis, 1 is again the identity operator.

Appendix C

Laplacian Operators

An important role in the definition of a convolution on simplicial complexes is played by the Laplacian operator, especially by its eigenfunctions and spectrum. A more thorough discussion of this operator can be found in [10].

Definition C.1. We define the p -Laplacian operator to be

$$\Delta_p := \partial_{p+1} \partial_{p+1}^\dagger + \partial_p^\dagger \partial_p =: \Delta_p^+ + \Delta_p^-.$$

The Laplacian operator is defined to be self-adjoint and positive definite.

Proposition C.2. Let Δ_p be a laplacian operator, then $\Delta_p^\dagger = \Delta_p$.

Proof. Let $|\sigma\rangle, |\tau\rangle \in C_p$

$$\begin{aligned} \langle \sigma | \Delta_p | \tau \rangle &= \langle \sigma | (\partial_{p+1} \partial_{p+1}^\dagger + \partial_p^\dagger \partial_p) | \tau \rangle = \\ &= \langle \tau | (\partial_{p+1} \partial_{p+1}^\dagger + \partial_p^\dagger \partial_p)^\dagger | \sigma \rangle = \\ &= \langle \tau | (\partial_{p+1} \partial_{p+1}^\dagger + \partial_p^\dagger \partial_p) | \sigma \rangle = \langle \tau | \Delta_p | \sigma \rangle. \end{aligned}$$

□

According to the spectral theorem there exists a basis of eigenchains of the Laplacian, and since all $\Delta_p, \Delta_p^+, \Delta_p^-$ are self-adjoint we can say that they all admit a basis of eigenchains.

Proposition C.3. Let $\Delta_p |\sigma\rangle = \lambda_\sigma |\sigma\rangle$ then $\lambda_\sigma \geq 0$.

Proof. Let $\Delta_p^+ |\sigma\rangle = \lambda_\sigma^+ |\sigma\rangle$, we see that $\langle \sigma | \Delta_p^+ | \sigma \rangle = |\partial_{p+1} |\sigma\rangle|^2 \geq 0$, and since $\langle \sigma | \Delta_p^+ | \sigma \rangle = \lambda_\sigma^+ \langle \sigma | \sigma \rangle$ we have that $\lambda_\sigma^+ \geq 0$.

Let then $\Delta_p^- |\sigma\rangle = \lambda_\sigma^- |\sigma\rangle$, we see that $\langle \sigma | \Delta_p^- | \sigma \rangle = |\partial_p |\sigma\rangle|^2 \geq 0$, and since $\langle \sigma | \Delta_p^- | \sigma \rangle = \lambda_\sigma^- \langle \sigma | \sigma \rangle$, we also have that $\lambda_\sigma^- \geq 0$.

Furthermore, since $\Delta_p^+ \Delta_p^- = \Delta_p^- \Delta_p^+ = 0$ we have that $[\Delta_p^+, \Delta_p^-] = 0$, thence $[\Delta_p, \Delta_p^\pm] = 0$, therefore $\Delta_p, \Delta_p^+, \Delta_p^-$ share a basis of eigenchains. Let $|\sigma\rangle$ be in that common basis then $\Delta_p |\sigma\rangle = \lambda_\sigma |\sigma\rangle$, where $\lambda_\sigma = \lambda_\sigma^+ + \lambda_\sigma^- \geq 0$. □

Another really interesting property that was first proven by Beno Eckmann in 1944, is that the kernel of the p -Laplacian is isomorphic to the p -homology group.

Theorem C.4. Let Δ_p be a laplacian operator, then $\ker \Delta_p \simeq H_p$.

Proof. We recall the definition of laplacian $\Delta_p =: \partial_{p+1} \partial_{p+1}^\dagger + \partial_p^\dagger \partial_p$.

Because of the homology lemma $\Delta_p^+ \Delta_p^- = \Delta_p^- \Delta_p^+ = 0$, therefore $\ker \Delta_p^\pm \subset \text{im } \Delta_p^\mp$.

It is trivial to see that, since $\ker \Delta_p = \ker \Delta_p^+ \cap \ker \Delta_p^-$, we have that $\ker \partial_{p+1}^\dagger \cap \ker \partial_p \subset \ker \Delta_p$.

Less trivial is the opposite inclusion, in fact, let $|\sigma\rangle \in \ker \Delta_p$, we have that $\partial_{p+1}^\dagger |\sigma\rangle \in \ker \partial_{p+1} =$

$(\text{im } \partial_{p+1}^\dagger)^\perp$ and of course that $\partial_{p+1}^\dagger |\sigma\rangle \in \text{im } \partial_{p+1}^\dagger$.

Therefore, since our integration is an inner product, because of the adjunction relation we have that the chain $\partial_{p+1}^\dagger |\sigma\rangle$ must be the null chain, hence the inclusion $\ker \Delta_p \subset \ker \partial_{p+1}^\dagger$.

Similarly we can see that $\ker \Delta_p \subset \ker \partial_p$, therefore $\ker \Delta_p \subset \ker \partial_{p+1}^\dagger \cap \ker \partial_p$. Now, since $\ker \Delta_p = \ker \partial_{p+1}^\dagger \cap \ker \partial_p$, we conclude that $\ker \Delta_p = (\text{im } \partial_{p+1}^\dagger)^\perp \cap \ker \partial_p \simeq H_p$. \square

Bibliography

- [1] J. A. Bondy and U. S. R. Murty. *Graph Theory with Applications*. Elsevier, 1976.
- [2] Michael M. Bronstein et al. “Geometric Deep Learning: Going beyond Euclidean data”. In: (2017).
- [3] Michael M. Bronstein et al. *Geometric Deep Learning: Grids, Groups, Graphs, Geodesics, and Gauges*. 2021. arXiv: 2104.13478 [cs.LG].
- [4] Taco S. Cohen and Max Welling. *Group Equivariant Convolutional Networks*. 2016.
- [5] S. Ebli, M. Defferrard, and G. Spreemann. *Simplicial Neural Networks*. 2020.
- [6] Herbert Edelsbrunner and John Harer. *Computational Topology - an Introduction*. American Mathematical Society, 2010.
- [7] I. M. Gelfand and S. V. Fomin. *Calculus of Variations*. Dover, 1963.
- [8] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. <http://www.deeplearningbook.org>. MIT Press, 2016.
- [9] Allen Hatcher. *Algebraic Topology*. Cambridge Univ. Press, 2000.
- [10] Danijela Horak and Jürgen Jost. *Spectra of combinatorial Laplace operators on simplicial complexes*. 2011.
- [11] Serge Lang. *Algebra*. Springer, New York, NY, 1955.
- [12] Saunders MacLane. *Categories for the Working Mathematician*. 1971.
- [13] Joseph J. Rotman. *An Introduction to Algebraic Topology*. Graduate Texts in Mathematics. Springer-Verlag New York, 1988.
- [14] I. M. Singer and J. A. Thorpe. *Lecture Notes on Elementary Topology and Geometry*. Springer, New York, NY, 1976.